

Curation Over Breadth: Design and Measurement of a Self-Hostable, Availability-Aware AI Gateway

Sai Vivek Peddi · Claude (benchmark harness & analysis)

MLPal · Technical report, August 2026 · measurements executed 11 August 2026 unless dated otherwise
Code, harness, and raw results: github.com/mlpalOld/mlpal-gateway (Apache-2.0)

Abstract

AI gateways compete on catalog breadth: OpenRouter advertises 400+ models, Portkey 1,600+, and LiteLLM 100+ providers [1,7,5]. We argue breadth is the wrong metric. Frontier models now retire on a ~ 12 -month cycle — Gemini 1.5 snapshots lived exactly 12 months, Claude Opus 4.1 was retired one year to the day after release, GPT-4.5-preview survived 4.5 months, and AWS Bedrock has made a 12-month platform lifetime *contractual* [9,11,12,14] — so a catalog's tail is mostly retired-or-dominated inventory, while the utility of a gateway concentrates in whether the handful of current models are served well and whether references to them survive churn. We present **MLPal Gateway**, an open-source (Apache-2.0), self-hostable gateway built on this hypothesis: a curated 55-entry catalog (73 models served in the managed deployment) behind an Anthropic-wire core with OpenAI-compatible endpoints, *availability-aware router tags* that resolve a stable alias (`mlpal`) to the best model a given deployment can actually serve, pass-through compute-unit metering, and per-key policy, budgets, and cache/latency observability. We measure it three ways. (i) **Gateway overhead**, isolated by co-locating client and gateways: against direct `api.anthropic.com` (median TTFT 510 ms, $N=15$), MLPal Gateway adds 38 ms at the median with a *tighter* tail than a minimally-configured LiteLLM proxy (p95 716 vs 754 ms; +12 ms median) — $\sim 1\%$ of a 3.1 s request in exchange for metering, policy, and budget admission. (ii) **Provider-semantic preservation**: prompt caching passes through byte-faithfully on a 22k-token prefix (write $\rightarrow$ read on all systems), and the gateway's metered cost reproduces Anthropic's list pricing to five decimals, falling $12.3\times$ on cache hits (0.002756 $\rightarrow$ 0.000224 CU, measured end-to-end). (iii) **System-level utility**: a coding agent using the gateway's ranked catalog to route sub-tasks cut sub-agent cost $\sim 10\times$ at near-matched correctness versus pinning a frontier model (dated 2026-07-28 [18]). A July-2026 head-to-head against OpenRouter from the same vantage point [17] is summarized as a prior study. We release the gateway, console, SDK, and this paper's full harness and raw data.

Keywords: LLM gateway · model routing · model deprecation · self-hosting · prompt caching · metering

1 Introduction

An AI gateway multiplexes many model providers behind one API and one billing relationship. The category has consolidated around a breadth race: OpenRouter's pricing page advertises 400+ models across 70+ providers [1]; Portkey's gateway README claims routing to 1,600+ models [7];

LiteLLM supports 100+ providers [5]. Model count is legible, easy to market, and — we argue — largely disconnected from the utility a gateway delivers.

Our hypothesis is that **provider agnosticism matters, but breadth does not**: models are retiring at a record and accelerating pace (§2), so most of a large catalog is dead or dominated weight, while real traffic concentrates on a small set of current models per provider. What a practitioner needs from a gateway is (a) that the current models are served *unhobbled* — correct parameter ranges, caching, streaming fidelity, provider quirks handled per model; (b) that their own code survives retirements without edits; and (c) that the cost and behavior of every request is observable. Each of these is a per-model investment that scales to tens of models, not thousands.

MLPal Gateway is a production gateway built on this thesis. This report contributes:

1. A quantified case that model lifetimes have compressed toward ~ 12 months, from providers' own deprecation ledgers (§2, Fig. 1).
2. The system design: an Anthropic-wire universal core with OpenAI-compatible endpoints, availability-aware router tags, pass-through compute-unit metering, and per-key admission (policy, budgets, rate limits) — one codebase serving both a managed deployment and a self-hosted open-source distribution (§3, Figs. 2–3).
3. Controlled measurements: gateway overhead isolated on localhost against direct provider access and a LiteLLM proxy (§5.1); managed-service latency from a client vantage point (§5.2); prompt-cache passthrough and metering fidelity verified end-to-end (§5.3); and system-level routing utility via a coding-agent workload (§5.4).
4. An honest feature comparison against OpenRouter, LiteLLM, Portkey, and Helicone from their current public documentation (§6), and the full reproduction harness.

2 The case for curation

Model lifetimes are compressing. Fig. 1 plots launch-to-retirement spans for twelve production API models, taken from providers' own lifecycle pages [9,11,12,14]. The 2024 cohort lived 22–29 months (claude-3-opus, gpt-4o's scheduled Azure retirement); the 2025 cohort converges on 12–16 months; GPT-4.5-preview — a flagship-priced model — survived **4.5 months** with a 91-day removal notice [9]. Google's numbered Gemini snapshots have retired exactly 12 months after release (1.5-pro-001: 2024-05-24→2025-05-24; -002: 2024-09-24→2025-09-24) [12]. Claude Opus 4.1 was retired one year to the day after release (2025-08-05→2026-08-05), with a ~ 61 -day notice [11,16].




Figure 1: Launch→retirement spans for production API models, ordered by launch date, from official provider lifecycle pages [9,11,12,14] (gpt-4o shows Microsoft's published Azure retirement schedule; Bedrock dates are platform-specific). The visible trend: models launched in 2024 lived 22–29 months; models launched in 2025 live 12–16. Notice windows compressed in parallel — Anthropic's 2025 retirements carried ~181-day notices, its 2026 cohort ~61 days [11,16]; OpenAI's policy floor is 6 months for GA models and 3 months for specialized variants [9].

Churn is now contractual. AWS Bedrock's lifecycle policy guarantees only 12 months of platform availability after launch, moves models to a Legacy state in which an account inactive on that model for 15 days loses access, and (for post-2026 EOLs) adds an extended-access phase at *provider-set higher pricing* [14]. Deprecation is no longer an exceptional event a gateway can ignore; it is a scheduled property of the substrate.

Retirement breaks production. When OpenAI removed GPT-4o from ChatGPT at GPT-5's launch with zero notice (2025-08-07), the user revolt forced reinstatement within five days [10,15]; the February-2026 retirement wave left some business customers without model access for a week [10]. On Claude Opus 4.1's retirement day, hard-coded model IDs failed simultaneously across Bedrock, Vertex, and the first-party API, and migration was complicated by a breaking parameter change in the successor [16]. OpenAI's own staff articulated the underlying tension: older models are kept in the API because "the alternative would be to risk breaking production apps whenever we launch a new model."

Serving a model well is per-model work. Getting a model's best behavior requires model-specific handling that a lowest-common-denominator passthrough cannot provide: reasoning-budget ranges that are discontinuous and differ per model (we maintain per-model valid ranges for Gemini's `thinking_budget`); prompt-cache minimums that differ per model (a ~2.5k-token prefix silently fails to cache on claude-haiku-4-5, whose minimum is 4,096 tokens — the request succeeds and simply bills full price); parameter support that changes across successor models [16]; and streaming/error quirks per provider. This work scales to a curated set, not to 1,600 models. A

catalog that is small, current, and tuned is therefore a feature, and the right unit of comparison is *utility per supported model*, not the count.

3 System design




Figure 2: Request path. Admission (auth → rate limit → billing gate → model policy → spend budgets) runs before routing; the router resolves either a concrete model or a router tag; adapters translate to each provider's native API behind circuit breakers and relay SSE without re-chunking. Metering, the usage ledger, cache/TTFT statistics, and optional payload capture run off the hot path after the response.

3.1 One codebase, two deployments

The managed service (`models.mlpal.ai`) and the open-source distribution (`mlpal-gateway`) are the same codebase. Deployment differences are confined to composition-root *seams* — auth backend (local keys vs. platform SSO), billing backend (local metering vs. platform wallet), and payload-capture default (on for self-hosters operating their own box; off in the managed service, which retains billing metrics but never payloads) — never forks. The distribution boots with one `docker compose up`: Postgres, Redis, the gateway, and an admin console; a catalog reconcile seeds the model registry, pricing, and routing feed; 341 unit tests and the console build gate CI.

3.2 Wire surface

The core endpoint is **Anthropic-wire** `POST /v1/messages` — request and response envelopes, SSE event grammar, tool use, and `cache_control` are Anthropic's, for every provider behind the gateway. OpenAI-compatible endpoints (`/v1/chat/completions`, `/v1/embeddings`, `/v1/images/generations`, `/v1/audio/*`) serve existing OpenAI-SDK code unchanged. A version prefix denotes a contract revision, never a wire dialect. Every response carries the request's metered cost in an `X-MLPal-Compute-Units` header and a routing envelope naming the resolved model.

3.3 Availability-aware router tags




Figure 3: Router-tag resolution. A tag (`mlpal`, `mlpal-flash`, `mlpal-lite`) maps to a priority-ordered, provider-spanning candidate list maintained in the catalog feed. Resolution walks the list and picks the first candidate the *local deployment* can serve (model active, not paused, provider key configured). The client-visible name never changes; retirements and upgrades are absorbed by feed updates.

Router tags operationalize the curation thesis. A client that ships `"model": "mlpal"` survives every retirement in Fig. 1 without a code change: when a candidate is retired or paused, resolution falls through to the next; when the feed promotes a new frontier model, the same tag starts serving it. Because candidate lists span providers, the mechanism degrades gracefully to single-provider deployments — an instance with only an Anthropic key still resolves all tags (measured in §5: the managed deployment currently resolves `mlpal` → `claude-opus-5`). The same curated intelligence is exposed declaratively as a ranked **catalog** (`GET /v1/catalog`, 55 entries with tier, capabilities, and per-token rates) for clients that choose their own model per task — the two consumption modes differ only in who decides, gateway or client (§5.4 measures the latter).

3.4 Metering: pass-through compute units

Cost is metered in *compute units* pegged at $1 \text{ CU} = \$10$ of provider list price: $\text{CU} = \sum_c \text{token-}s_c \cdot \text{rate}_c / 10$ over billing components c (input, output, cache-write at the provider's $1.25\times$ input rate, cache-read at $0.1\times$). There is no markup term: self-hosters pay providers directly, and the managed tier bills tokens at cost with a flat platform fee, so the meter is an *instrument*, not a margin. The ledger stores CU at `numeric(24,12)` — sub-nano-dollar requests meter without truncation — and §5.3 verifies the meter reproduces provider list pricing end-to-end. Per-key spend budgets (multiple sliding windows, checked at admission, never cutting a running stream) and model-policy globs (allow/deny, evaluated after alias resolution) complete the admission pipeline.

3.5 Observability

Every request writes a usage row (trace ID, resolved model, tokens, CU, latency, stream flag, time-to-first-token, cache read/write tokens); per-key aggregates expose cache hit rate, latency p50/p95, TTFT p50, and stream share; payload capture (request/response bodies, zlib-compressed, runtime-toggable, opt-in) is a self-hosting feature — the operator owns the box and the data. The managed service keeps the metrics and never stores payloads.

4 Experimental setup

Systems. Four systems serve the same model snapshot (`claude-haiku-4-5-20251001`) over the same upstream: (a) **direct** — `api.anthropic.com/v1/messages`; (b) **MLPal Gateway** — this work, Docker on localhost, full admission enabled (auth, Redis rate limiting, billing gate, policy, budgets, metering, usage persistence); (c) **LiteLLM proxy** — `ghcr.io/berriai/litellm:main-latest` (pulled 2026-08-11), Docker on localhost, master-key auth, its native Anthropic `/v1/messages` endpoint [6]; (d) **MLPal managed** — `models.mlpal.ai` (AWS us-east-2, internet-facing ALB), reached over a residential connection, so (d) includes real WAN/TLS costs that (a–c) partially avoid. Systems (b) and (c) call the provider with the same API key from the same machine, so their deltas against (a) isolate **gateway processing overhead**: admission, translation, relay, and metering. LiteLLM is configured minimally (no logging callbacks, no database); it represents a thin-proxy floor rather than a tuned production deployment.

Protocol. Identical prompt (a $\sim$ 150-word CAP-theorem explanation task), `max_tokens` 256, streaming SSE, one cold client connection per request, one unmeasured warmup per system, then $N=15$ measured trials per system with A/B/C order rotated per trial to cancel temporal drift; 0.5 s between requests. **Metrics.** TTFB (first SSE byte), TTFT (first `content_block_delta`), total latency, and content-chunk count, from a single client (Apple-silicon macOS, Python httpx). **Cache experiment.** A $\sim$ 22k-token system prefix marked `cache_control: ephemeral` (well above haiku-4.5's 4,096-token cache minimum), distinct per system so systems cannot share provider cache entries; 4 sequential non-streaming trials per system; usage fields read from the wire. Metering fidelity is verified on system (b) by reading the CU header on a cache-write and a cache-read call of the same prefix.

5 Results

5.1 Gateway overhead




Figure 4: Per-trial TTFT and total latency (N=15 per system, streaming, cold connections; vertical bar = median). The three localhost systems are within ~ 40 ms of each other at the median — indistinguishable relative to a ~ 3.1 s request. The managed row is a different experiment (adds real WAN+TLS from a residential vantage) and is shown for context, not as an overhead measurement.

Table 1: Streaming results, N=15 measured trials per system. Δ median is TTFT median minus direct's. MLPal Gateway runs full admission (rate limit, billing gate, policy, budgets, metering, persistence); LiteLLM runs master-key auth only.

System	TTFT med (ms)	mean $\pm$ std	p95	Δ med	Total med (ms)	Chunks med
Direct (api.anthropic.com)	510	517 $\pm$ 68	634	—	3,077	9
MLPal Gateway (localhost)	548	568 $\pm$ 80	716	+38	3,180	9
LiteLLM proxy (localhost)	522	557 $\pm$ 112	754	+12	3,015	9
MLPal managed (WAN, us-east-2)	898	1,017 $\pm$ 454	1,848	+388 [†]	3,523	10

[†] includes residential-network RTT and per-request TLS setup to us-east-2; not a gateway-overhead measurement.

Three observations. First, **gateway overhead is noise at the median**: +38 ms (MLPal) and +12 ms (LiteLLM) against a 3,077 ms total — 1.2% and 0.4% of the request. The ~ 25 ms between the two gateways buys the full admission pipeline (per-key policy, sliding-window budgets, CU metering, usage persistence) versus master-key auth. Second, **the tail ordering inverts**: MLPal's TTFT dispersion is tighter than LiteLLM's (std 80 vs 112 ms; p95 716 vs 754 ms) despite the higher median — overhead that is constant is cheaper at the tail than overhead that is variable. Third, **stream fidelity is preserved**: all three systems deliver the same median 9 content chunks; neither gateway coalesces or buffers the provider's stream.

5.2 The managed service, and the July head-to-head with OpenRouter

From this vantage point the managed deployment serves haiku at 898 ms median TTFT — consistent with the 922 ms measured in our July study from the same client [17], i.e., the /v1 surface migration and admission additions since July did not regress client-visible latency. Fresh verification of the surface (2026-08-11): 73 models served across four providers (37 OpenAI, 14

Anthropic, 14 Google, 8 Bedrock — the latter curating current open-weight models: GLM-5, Kimi K2.5, DeepSeek V3.2, Qwen3-Coder-480B, Llama-4-Maverick, gpt-oss-120b), a 55-entry curated catalog, `mlpal` → `claude-opus-5` resolution, and a CU header on every response (0.0000145 CU for a 16-output-token haiku call).

The July-2026 study [17] measured 11 matched model pairs against OpenRouter (110 paired streaming requests plus supplements) from this same vantage. Summary, quoted as a dated prior result: despite a structural ~ 190 ms connection-setup handicap (single-region ALB vs. CDN edge), MLPal won median TTFT on 8 of 11 pairs — including every Anthropic model tested — with substantially tighter tails (p95 within $1.1\text{--}2.6\times$ of median on 10 of 11 models, where OpenRouter's multi-provider load balancing produced $4.3\text{--}6.2\times$ blowups on several); OpenRouter won median TTFT on fast/cheap models where edge proximity dominates, sustained 10–30% higher steady-stream throughput, and finished streamed responses sooner on 8 of 11 pairs; OpenRouter returned 3 empty completions (reasoning consumed the token budget) and exhibited fully-buffered streaming on some routes (Gemini Flash arriving in 1–2 chunks vs. MLPal's ~ 9), while MLPal returned visible content 67/67 times. Neither gateway dominated; the trade is breadth and resilience against consistency and stream fidelity. Those numbers predate this report's changes and OpenRouter's own evolution since July; we cite rather than restate them.

5.3 Provider semantics survive the hop: caching and metering fidelity




Figure 5: Prompt-cache passthrough, $\sim 22\text{k}$ -token prefix with `cache_control`, 4 sequential trials per system (distinct prefixes per system). Every system shows the canonical pattern — full cache write on trial 1, full cache read on trials 2–4 — confirming both gateways relay Anthropic's caching semantics byte-faithfully.

Prompt caching is the highest-leverage provider feature a gateway can break silently: a gateway that reorders or rewrites the prefix destroys cache identity and the client pays full price with no error. All three systems preserved it exactly (write 21,213–22,013 tokens on trial 1; read of the identical count on trials 2–4). On MLPal, metering tracked it end-to-end: the cache-write call metered **0.002756 CU** and the cache-read call **0.000224 CU** — a **12.3 $\times$** reduction, and each figure reproduces Anthropic's list price exactly ($22,013 \text{ tokens} \times 1.25 \times \$1/\text{M input} \div \$10/\text{CU} = 0.00275$; $\times 0.1$ for reads = 0.00022) — i.e., the meter is faithful to five decimal places, verified from the wire rather than from our own configuration. Cache read/write tokens land in each usage row, so

per-key cache hit rate is a first-class dashboard metric; on our development deployment, an agentic coding workload (yodex) sustains a 44% per-key cache hit rate that is directly visible there.

5.4 Utility over breadth, measured at the system level

The curation thesis predicts that a *ranked, current* catalog beats a large flat one for real workloads. A measurement from our coding agent (yodex, 2026-07-28 [18]) exercises exactly this: on a 4-module build with delegation, the agent consulted the gateway catalog to route each sub-task by rated complexity — trivial modules to `gemini-2.5-flash-lite`, harder ones upward — cutting sub-agent cost $\sim 10\times$ ($0.212 \rightarrow 0.022$ CU) versus inheriting the pinned frontier model, at near-matched correctness (11/12 vs 12/12), and $\sim 2\times$ cheaper end-to-end than a same-model Claude Code baseline. The enabling primitive is not 1,600 models; it is a catalog small enough to rank honestly and current enough to trust.

6 Feature comparison

Table 2: Feature surfaces from current public documentation (retrieved 2026-08-11) [1,2,4,5,6,7,8] and MLPal Gateway source [3]. "—" = not offered/documented; *italics* = the vendor's hosted/enterprise tier, not the OSS artifact.

Dimension	OpenRouter	LiteLLM (OSS)	Portkey (OSS)	Helicone	MLPal Gateway
Catalog	400+ models, 70+ providers	100+ providers	1,600+ models claimed	100+ models	55 curated entries; 73 served (managed)
Thesis	breadth + multi-provider resilience	universal adapter	breadth + guardrails	observability-first	curated, current, unhobbled
Self-host	—	MIT	MIT	Apache-2.0 (cloud-primary posture)	Apache-2.0 core (+ commercial dir)
Anthropic-wire endpoint	— (OpenAI wire + Responses beta)	yes	not documented	not documented	core surface, all providers
Fee model	5.5% on credits; BYOK tiers	free / quote enterprise	free / \$49+ hosted logs	0% + Stripe passthrough	free self-host; managed: tokens at cost + flat fee
Stable model alias	<code>openrouter/auto</code> (auto-router)	alias config (operator-defined)	configs/routing rules	—	router tags: curated, availability-aware, feed-updated
Multi-provider fallback (same model)	yes, price-weighted	yes (strategies)	yes	yes	— (one pinned first-party edge per model)
Per-request cost on wire	via /generation lookup	response header (cost tracking)	<i>hosted dashboards</i>	registry-computed	CU header + routing envelope, every response
Per-key cache/TTFT/tail metrics	activity dashboard	via callbacks (Langfuse etc.)	<i>hosted</i>	core product	built-in per key: cache hit rate, TTFT p50, p50/p95
Spend controls	credit limits per key	budgets per key/team	<i>hosted budgets</i>	tier limits	multi-window budgets + model-policy globs,

Dimension	OpenRouter	LiteLLM (OSS)	Portkey (OSS)	Helicone	MLPal Gateway
					admission-time
Payload capture (own box)	n/a (hosted)	log callbacks	local console logs	core product (hosted)	opt-in, zlib, runtime toggle; managed never stores payloads

The surfaces reflect each system's thesis, and none dominates. OpenRouter's multi-provider failover, auto-router, and compliance controls have no MLPal equivalent; LiteLLM's provider matrix is far larger; Helicone's observability product is deeper than our console. MLPal's differences are the Anthropic-wire universal core, availability-aware curated tags, admission-time policy/budget enforcement, and a meter whose output we can verify against provider list price on the wire (§5.3) — investments that only make sense on a curated catalog.

7 Limitations and conflict of interest

All measurements come from a single client machine on a single day; N=15 per cell makes medians robust but p95 estimates coarse; one prompt shape, one model family (haiku-4.5) for the overhead study; totals include provider-side generation variance that dwarfs gateway deltas. The local-host design isolates gateway *processing* overhead but not deployment topology — a production LiteLLM with logging callbacks and a database, or MLPal behind a load balancer, will differ. LiteLLM was measured in its minimal configuration; its overhead should be read as a thin-proxy floor, not a verdict on the project. The managed-service row includes network costs the localhost rows do not. The July OpenRouter study is quoted as-dated; OpenRouter has shipped changes since. Catalog counts across vendors are not directly comparable (models vs providers vs integrations). Output quality was not evaluated — only delivery, semantics preservation, and cost. This report is written by the system's authors; the harness, raw JSON, and figure code ship in the repository so every number can be re-derived, and the strongest competitive claims (§5.2) are the ones we measured *against* ourselves.

8 Conclusion

Breadth is the legible metric, but the substrate has changed under it: when models live ~12 months and churn is contractual, a catalog's value decays at the rate its entries retire, while the gateway properties that compound — stable aliases that survive retirement, per-model tuning, faithful caching, verifiable metering, admission-time governance — are all per-model investments that favor curation. MLPal Gateway implements this position and measures its costs honestly: ~38 ms of median overhead for a full admission pipeline, tails tighter than a thin proxy, provider se-

mantics preserved to the token, and a meter faithful to five decimals. The gateway, console, SDK, and this paper's harness are open source (Apache-2.0) at github.com/mlpalOld/mlpal-gateway; the managed deployment serves the same code at models.mlpal.ai.

References

1. OpenRouter, "Pricing." openrouter.ai/pricing (retrieved 2026-08-11).
2. OpenRouter, "Provider selection & routing." openrouter.ai/docs/guides/routing/provider-selection.
3. MLPal, "mlpal-gateway" source. github.com/mlpalOld/mlpal-gateway.
4. OpenRouter, "Zero completion insurance." openrouter.ai/docs/guides/features/zero-completion-insurance (retrieved 2026-08-11).
5. LiteLLM, project and proxy docs: load balancing & routing strategies. [litellm.ai](https://litellm.ai/docs); docs.litellm.ai/docs/proxy/load_balancing.
6. LiteLLM, "/v1/messages (Anthropic unified endpoint)." docs.litellm.ai/docs/anthropic_unified.
7. Portkey, "AI Gateway" README (MIT; "1,600+ language, vision, audio, and image models"). github.com/portkey-ai/gateway (retrieved 2026-08-11).
8. Helicone, "AI Gateway" and credits (0% markup passthrough). docs.helicone.ai/gateway/overview; helicone.ai/credits.
9. OpenAI, "Deprecations." developers.openai.com/api/docs/deprecations (retrieved 2026-08-11): GA ≥6-month / specialized ≥3-month policy; gpt-4.5-preview announced 2025-04-14, shutdown 2025-07-14; o1-preview 2025-04-28→2025-07-28.
10. OpenAI, "Retiring GPT-4o and older models" (2026-02-13). openai.com/index/retiring-gpt-4o-and-older-models; TechRadar, "OpenAI pulls older ChatGPT models..." (2026-02).
11. Anthropic, "Model deprecations." platform.claude.com/docs/en/about-claude/model-deprecations (retrieved 2026-08-11): ≥60-day policy; Claude 1.x retired 2024-11-06 (63-day notice); Claude 2.x/3-Sonnet retired 2025-07-21 (181-day); Claude 3 Opus retired 2026-01-05.
12. Google Cloud, "Model versions and lifecycle." docs.cloud.google.com/gemini-enterprise-agent-platform/models/model-versions: gemini-1.5-pro-001 2024-05-24→2025-05-24; -002 2024-09-24→2025-09-24; gemini-2.0-flash 2025-02-05→2026-06-01. PaLM API decommissioned 2024-08-15 (ai.google.dev/palm_docs/deprecation).
13. Microsoft Azure, "Model retirement schedule." learn.microsoft.com/azure/foundry/openai/concepts/model-retirement-schedule (gpt-4o: 2024-05-13 launch, 2026-10-01 scheduled retirement).
14. AWS, "Amazon Bedrock model lifecycle." docs.aws.amazon.com/bedrock/latest/userguide/model-lifecycle.html: ≥12-month platform availability, ≥6-month Legacy, 15-day-inactivity access loss, post-2026 extended-access pricing; Llama-3.1-405B EOL 2026-07-07.
15. Ars Technica, "OpenAI brings back GPT-4o after user revolt" (2025-08). arstechnica.com.
16. TheRouter, "Anthropic model deprecation cadence" (third-party analysis; 8 Claude retirements in 12 months to 2026-05; 60–62-day notices; Opus-4.1 2025-08-05→2026-08-05). therouter.ai/blog/anthropic-model-deprecation-cadence-migration-guide. Third-party tracker; key dates cross-checked against [11].
17. Peddi & Claude, "When the Edge Isn't Enough: A Head-to-Head Latency and Capability Study of Two LLM Gateways." MLPal technical report, 2026-07-18. Benchmark executed from the same client vantage point as this

report.

18. MLPal, "Yodex vs Claude Code on claude-opus-5" (2026-07-28). github.com/mlpal0ld/yodex (benchmarks/).
19. Anthropic, "Prompt caching" (cache-write $1.25\times$ input rate, cache-read $0.1\times$; per-model minimum cacheable prefix). platform.claude.com/docs/en/build-with-claude/prompt-caching.

Appendix A Reproduction

`paper/bench/` in the repository contains the harness (`harness.py`, `cache_exp.py`), the LiteLLM configuration used, raw results (`results-*.json`), figure code, and aggregate statistics (`stats.py`). Environment: set `ANTHROPIC_API_KEY`, a gateway key (`MLPAL_KEY`), boot the gateway with `docker compose up`, start LiteLLM with the shipped config, then `python harness.py 15 direct,mlpal,litellm results.json`. The managed-service rows require only a key. Model snapshot: `claude-haiku-4-5-20251001`; client: macOS/Apple silicon, Python 3.13 + httpx, residential network, 2026-08-11.